

Predict the survival of passengers based on features.

- Features: age, fare, sex, and others.
- Classes: 2 (Survived or Not Survived).

## Logistic Regression

```
In [1]: #Import necessary Libraries

import pandas as pd
import numpy as np
import mglearn
import matplotlib.pyplot as plt
from pandas.plotting import scatter_matrix
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [2]: #Load the dataset

url = "https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv"
titanic = pd.read_csv(url)
```

```
In [3]: #Preprocess the data

titanic = titanic.drop(['PassengerId', 'Name', 'Ticket', 'Cabin'], axis=1)
titanic['Age'].fillna(titanic['Age'].mean(), inplace=True)
titanic['Embarked'].fillna(titanic['Embarked'].mode()[0], inplace=True)
```

```
In [4]: titanic.head()
```

```
Out[4]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	male	22.0	1	0	7.2500	S
1	1	1	female	38.0	1	0	71.2833	C
2	1	3	female	26.0	0	0	7.9250	S
3	1	1	female	35.0	1	0	53.1000	S
4	0	3	male	35.0	0	0	8.0500	S

```
In [5]: #Convert categorical variables into dummy/indicator variables

titanic = pd.get_dummies(titanic, columns=['Sex', 'Embarked'], drop_first=True)
```

```
In [6]: titanic.head()
```

Out[6]:

	Survived	Pclass	Age	SibSp	Parch	Fare	Sex_male	Embarked_Q	Embarked_S
0	0	3	22.0	1	0	7.2500	1	0	1
1	1	1	38.0	1	0	71.2833	0	0	0
2	1	3	26.0	0	0	7.9250	0	0	1
3	1	1	35.0	1	0	53.1000	0	0	1
4	0	3	35.0	0	0	8.0500	1	0	1

In [7]:

```
#Define features (X) and target (y)

X = titanic.drop('Survived', axis=1)
y = titanic['Survived']
```

In [8]:

```
#Split data into training and test sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=4)
```


In [9]:

```
print(X_train.shape)
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)
```

```
(712, 8)
(712,)
(179, 8)
(179,)
```

In [10]:

```
grr = scatter_matrix(X_train,
                      c=y_train,
                      figsize=(15,15),
                      marker='o',
                      s=60,
                      hist_kwds={'bins':20},
                      alpha=0.8,
                      cmap=mglearn.cm3)
```



In [11]:

```
#Train the Logistic Regression model
```

```
logreg = LogisticRegression(solver = 'liblinear')
logreg.fit(X_train, y_train)
```

Out[11]: LogisticRegression(solver='liblinear')

In [12]:

```
#Make predictions on the test set
```

```
y_pred = logreg.predict(X_test)
```

In [13]:

```
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)
```

```
In [14]: print(f"Accuracy: {accuracy}")
print("Confusion Matrix:")
print(conf_matrix)
print("Classification Report:")
print(class_report)
```

```
Accuracy: 0.7821229050279329
Confusion Matrix:
[[89 16]
 [23 51]]
Classification Report:
              precision    recall  f1-score   support

     0       0.79       0.85       0.82       105
     1       0.76       0.69       0.72       74

 accuracy
macro avg       0.78       0.77       0.77       179
weighted avg       0.78       0.78       0.78       179
```

## Logistic Regression (K-Fold)

```
In [15]: from sklearn.model_selection import cross_val_score, KFold
```

```
In [16]: #Perform K-Fold Cross-Validation

kf = KFold(n_splits=5, shuffle=True, random_state=42) # 5 folds
cv_scores = cross_val_score(logreg, X, y, cv=kf, scoring='accuracy')
```

```
In [17]: #Print the results of cross-validation

formatted_scores = [f"{score:.3f}" for score in cv_scores]
print(f"Cross-Validation Accuracy Scores: {formatted_scores}")

print(f"Mean Cross-Validation Accuracy: {cv_scores.mean():.3f}")
print(f"Standard Deviation of CV Accuracy: {cv_scores.std():.3f}")

Cross-Validation Accuracy Scores: ['0.782', '0.787', '0.848', '0.764', '0.815']
Mean Cross-Validation Accuracy: 0.799
Standard Deviation of CV Accuracy: 0.029
```

```
In [18]: logreg.fit(X, y)
y_pred = logreg.predict(X)
```

```
In [19]: accuracy = accuracy_score(y, y_pred)
conf_matrix = confusion_matrix(y, y_pred)
class_report = classification_report(y, y_pred)
```

```
In [20]: print(f"\nAccuracy on the whole dataset: {accuracy}")
print("Confusion Matrix:")
print(conf_matrix)
```

```
print("Classification Report:")
print(class_report)
```

Accuracy on the whole dataset: 0.8024691358024691

Confusion Matrix:

```
[[479  70]
 [106 236]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.87	0.84	549
1	0.77	0.69	0.73	342
accuracy			0.80	891
macro avg	0.80	0.78	0.79	891
weighted avg	0.80	0.80	0.80	891

## Logistic Regression (Grid Search)

```
In [21]: from sklearn.model_selection import GridSearchCV
```

```
In [22]: #Define the parameter grid

param_grid = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100],
    'penalty': ['l1', 'l2', 'elasticnet', 'none'],
    'solver': ['liblinear', 'saga']
}
```

```
In [23]: #Use GridSearchCV without K-Fold cross-validation

grid_search = GridSearchCV(logreg, param_grid, scoring='accuracy', verbose=1, n_jobs=-1)
grid_search.fit(X_train, y_train)
```

Fitting 5 folds for each of 48 candidates, totalling 240 fits

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_search.py:918: UserWarning: One or more of the test scores are non-finite: [0.66710332 0.66570472 0.67971043
0.68530484          nan          nan
          nan 0.68109918 0.66704422 0.66704422 0.7106274 0.68109918
          nan          nan          nan 0.68109918 0.7850586 0.68391608
0.7892938 0.67969073          nan          nan          nan 0.68109918
0.79630651 0.68109918 0.79629666 0.68109918          nan          nan
          nan 0.67969073 0.78925441 0.68109918 0.79207131 0.68109918
          nan          nan          nan 0.68109918 0.78925441 0.68109918
0.78925441 0.68109918          nan          nan          nan 0.68109918]
warnings.warn(
```

```
Out[23]: GridSearchCV(estimator=LogisticRegression(solver='liblinear'), n_jobs=-1,
    param_grid={'C': [0.001, 0.01, 0.1, 1, 10, 100],
    'penalty': ['l1', 'l2', 'elasticnet', 'none'],
    'solver': ['liblinear', 'saga']},
    scoring='accuracy', verbose=1)
```

```
In [24]: #Get the best parameters and best score from the grid search

best_params = grid_search.best_params_
best_score = grid_search.best_score_
```

```
print(f"Best Parameters from Grid Search: {best_params}")
print(f"Best Training Accuracy: {best_score:.3f}")
```

Best Parameters from Grid Search: {'C': 1, 'penalty': 'l1', 'solver': 'liblinear'}  
Best Training Accuracy: 0.796

In [25]:

```
#Evaluate on the test set
```

```
y_pred = grid_search.predict(X_test)
test_accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)
```

In [26]:

```
print(f"\nAccuracy on the test set: {test_accuracy:.3f}")
print("Confusion Matrix:")
print(conf_matrix)
print("Classification Report:")
print(class_report)
```

Accuracy on the test set: 0.793

Confusion Matrix:

```
[[88 17]
 [20 54]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.84	0.83	105
1	0.76	0.73	0.74	74
accuracy			0.79	179
macro avg	0.79	0.78	0.79	179
weighted avg	0.79	0.79	0.79	179

## Logistic Regression (K-Fold & Grid Search)

In [27]:

```
#Define the parameter grid
```

```
param_grid = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100],
    'penalty': ['l1', 'l2', 'elasticnet', 'none'],
    'solver': ['liblinear', 'saga']
}
```

In [28]:

```
#Set up K-Fold cross-validation and GridSearchCV
```

```
kf = KFold(n_splits=5, shuffle=True, random_state=42)
```

In [29]:

```
#Use GridSearchCV to perform the grid search
```

```
grid_search = GridSearchCV(logreg, param_grid, cv=kf, scoring='accuracy', verbose=1, n_
grid_search.fit(X, y)
```

Fitting 5 folds for each of 48 candidates, totalling 240 fits

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_search.py:918: UserWarning: One or more of the test scores are non-finite: [0.66220576 0.65883498 0.6868872
0.68351641 nan nan
nan 0.6868872 0.67004582 0.66555772 0.71267968 0.6868872
nan nan nan 0.6868872 0.79124976 0.68576988
0.79687402 0.6868872 nan nan nan 0.6868872
0.79573787 0.6868872 0.79912121 0.6868872 nan nan
nan 0.6868872 0.79908982 0.6868872 0.80471408 0.6868872
nan nan nan 0.6868872 0.80021342 0.6868872
0.80021342 0.6868872 nan nan nan 0.6868872 ]
warnings.warn(
```

```
Out[29]: GridSearchCV(cv=KFold(n_splits=5, random_state=42, shuffle=True),
estimator=LogisticRegression(solver='liblinear'), n_jobs=-1,
param_grid={'C': [0.001, 0.01, 0.1, 1, 10, 100],
'penalty': ['l1', 'l2', 'elasticnet', 'none'],
'solver': ['liblinear', 'saga']},
scoring='accuracy', verbose=1)
```

```
In [30]: #Get the best parameters and best score from the grid search
```

```
best_params = grid_search.best_params_
best_score = grid_search.best_score_

print(f"Best Parameters from Grid Search: {best_params}")
print(f"Best Cross-Validation Accuracy: {best_score:.3f}")
```

```
Best Parameters from Grid Search: {'C': 10, 'penalty': 'l2', 'solver': 'liblinear'}
Best Cross-Validation Accuracy: 0.805
```

```
In [31]: #Train the model with the best parameters
```

```
best_logreg = grid_search.best_estimator_
y_pred = best_logreg.predict(X)
```

```
In [32]: #Evaluate the model on the entire dataset
```

```
accuracy = accuracy_score(y, y_pred)
conf_matrix = confusion_matrix(y, y_pred)
class_report = classification_report(y, y_pred)
```

```
In [33]: print(f"\nAccuracy on the entire dataset: {accuracy:.3f}")
print("Confusion Matrix:")
print(conf_matrix)
print("Classification Report:")
print(class_report)
```

```
Accuracy on the entire dataset: 0.799
```

```
Confusion Matrix:
```

```
[[473  76]
 [103 239]]
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.82	0.86	0.84	549
1	0.76	0.70	0.73	342
accuracy			0.80	891
macro avg	0.79	0.78	0.78	891

weighted avg      0.80      0.80      0.80      891

## KNN

```
In [34]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=5)
```

```
In [35]: knn.fit(X_train, y_train)
```

```
Out[35]: KNeighborsClassifier()
```

```
In [36]: y_pred = knn.predict(X_test)
```

```
In [37]: print("Test set score: {:.2f}".format(knn.score(X_test, y_test)))

accuracy = accuracy_score(y_test, y_pred)
print("Model accuracy: {:.2f}" .format(accuracy))
```

Test set score: 0.69  
Model accuracy: 0.69

## Summary

```
In [41]: summary = {
    'Logistic Regression' : 0.782,
    'Logistic with K-Fold' : 0.802,
    'Logistic with Grid Serch' : 0.782,
    'Logistic with K-Fold & Grid Search' : 0.799,
    'KNN' : 0.71
}
```

```
In [42]: summary_df = pd.DataFrame(list(summary.items()), columns=['Model', 'Accuracy'])
```

```
In [43]: summary_df
```

```
Out[43]:
```

	Model	Accuracy
0	Logistic Regression	0.782
1	Logistic with K-Fold	0.802
2	Logistic with Grid Serch	0.782
3	Logistic with K-Fold & Grid Search	0.799
4	KNN	0.710