

K-Fold Cross Validation

```
In [1]: from sklearn.model_selection import cross_val_score
from sklearn.datasets import load_iris
from sklearn.linear_model import LogisticRegression
import mglearn
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: iris = load_iris()
logreg = LogisticRegression()
```

```
In [3]: scores = cross_val_score(logreg, iris.data, iris.target)
print("Cross-validation scores: {}".format([f"{score:.3f}" for score in scores]))
```

Cross-validation scores: ['0.967', '1.000', '0.933', '0.967', '1.000']

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

By default, cross_val_score performs 5-fold cross-validation, returning five accuracy values.

```
In [4]: scores = cross_val_score(logreg, iris.data, iris.target, cv=3)
print("Cross-validation scores: {}".format([f"{score:.3f}" for score in scores]))
```

Cross-validation scores: ['0.980', '0.960', '0.980']

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression


n_iter_i = _check_optimize_result(

```
In [5]: print("Average cross-validation score: {:.2f}".format(scores.mean()))
```

Average cross-validation score: 0.97

Using the mean cross-validation we can conclude that we expect the model to be around 97% accurate on average.

```
In [6]: mglearn.plots.plot_stratified_cross_validation()
```



It is usually a good idea to use stratified k-fold cross-validation instead of k-fold cross-validation to evaluate a classifier, because it results in more reliable estimates of generalization performance

More control over cross-validation

```
In [7]: from sklearn.model_selection import KFold
```

```
In [8]: kfold = KFold(n_splits=3, shuffle=True, random_state=0)

print("Cross-validation scores:\n{}".format(
    cross_val_score(logreg, iris.data, iris.target, cv=kfold)))
```

Cross-validation scores:
[0.98 0.96 0.96]

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

Leave-one-out Cross Validation

```
In [9]: from sklearn.model_selection import LeaveOneOut
```

```
In [10]: loo = LeaveOneOut()
scores = cross_val_score(logreg, iris.data, iris.target, cv=loo)
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression


```
n_iter_i = _check_optimize_result(
```

```
In [11]: print("Mean accuracy: {:.2f}".format(scores.mean()))
```

Mean accuracy: 0.97

Shuffle-split Cross-Validation

```
In [12]: mglearn.plots.plot_shuffle_split()
```



```
In [13]: from sklearn.model_selection import ShuffleSplit
```

```
In [14]: shuffle_split = ShuffleSplit(test_size=.8, train_size=.2, n_splits=10)
scores = cross_val_score(logreg, iris.data, iris.target, cv=shuffle_split)
```

```
In [15]: print("Cross-validation scores:\n{}".format([f"{score:.3f}" for score in scores]))
```

Cross-validation scores:

```
['0.942', '0.933', '0.950', '0.950', '0.942', '0.875', '0.933', '0.900', '0.925', '0.983']
```

Grid Search

```
In [16]: from sklearn.svm import SVC
         from sklearn.model_selection import train_test_split
```

```
In [17]: X_train, X_test, y_train, y_test = train_test_split(
         iris.data, iris.target, random_state=0)
```

```
In [18]: best_score = 0

         for gamma in [0.001, 0.01, 0.1, 1, 10, 100]:
             for C in [0.001, 0.01, 0.1, 1, 10, 100]:

                 svm = SVC(gamma=gamma, C=C)
                 svm.fit(X_train, y_train)

                 score = svm.score(X_test, y_test)

                 # if we got a better score, store the score and parameters
                 if score > best_score:
                     best_score = score
                     best_parameters = {'C': C, 'gamma': gamma}
```

```
In [19]: svm = SVC(**best_parameters)
         svm.fit(X_train, y_train)
         test_score = svm.score(X_test, y_test)
```

```
In [20]: print("Best score on validation set: {:.2f}".format(best_score))
         print("Best parameters: ", best_parameters)
         print("Test set score with best parameters: {:.2f}".format(test_score))
```

```
Best score on validation set: 0.97
Best parameters: {'C': 100, 'gamma': 0.001}
Test set score with best parameters: 0.97
```

Grid Search with Cross-Validation (1)

```
In [21]: for gamma in [0.001, 0.01, 0.1, 1, 10, 100]:
         for C in [0.001, 0.01, 0.1, 1, 10, 100]:

             svm = SVC(gamma=gamma, C=C)

             scores = cross_val_score(svm, X_train, y_train, cv=5)
             score = np.mean(scores)

             if score > best_score:
                 best_score = score
                 best_parameters = {'C': C, 'gamma': gamma}
```

```
In [22]: svm = SVC(**best_parameters)
svm.fit(X_train, y_train)
test_score = svm.score(X_test, y_test)
```

```
In [23]: print("Best score on validation set: {:.2f}".format(best_score))
print("Best parameters: ", best_parameters)
print("Test set score with best parameters: {:.2f}".format(test_score))
```

Best score on validation set: 0.97
 Best parameters: {'C': 100, 'gamma': 0.001}
 Test set score with best parameters: 0.97

Grid Search with Cross-Validation (2)

```
In [24]: from sklearn.model_selection import GridSearchCV
```

```
In [25]: param_grid = {'C': [0.001, 0.01, 0.1, 1, 10, 100],
                      'gamma': [0.001, 0.01, 0.1, 1, 10, 100]}

grid_search = GridSearchCV(SVC(), param_grid, cv=5)
```

```
In [26]: X_train, X_test, y_train, y_test = train_test_split(
iris.data, iris.target, random_state=0)
```

```
In [27]: grid_search.fit(X_train, y_train)
```

```
Out[27]: GridSearchCV(cv=5, estimator=SVC(),
                      param_grid={'C': [0.001, 0.01, 0.1, 1, 10, 100],
                                'gamma': [0.001, 0.01, 0.1, 1, 10, 100]})
```

```
In [28]: print("Test set score: {:.2f}".format(grid_search.score(X_test, y_test)))
```

Test set score: 0.97

```
In [29]: print("Best parameters: {}".format(grid_search.best_params_))
print("Best cross-validation score: {:.2f}".format(grid_search.best_score_))
```

Best parameters: {'C': 10, 'gamma': 0.1}
 Best cross-validation score: 0.97

Analyzing the result of cross-validation

```
In [30]: results = pd.DataFrame(grid_search.cv_results_)
results.head(3)
```

```
Out[30]:
```


	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_gamma	params	split
0	0.001595	0.000798	0.000405	0.000497	0.001	0.001	{'C': 0.001,	

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_gamma	params	split
							'gamma': 0.001}	
1	0.000797	0.000399	0.000599	0.000489	0.001	0.01	{'C': 0.001, 'gamma': 0.01}	
2	0.000792	0.000396	0.000598	0.000488	0.001	0.1	{'C': 0.001, 'gamma': 0.1}	

```
In [31]: scores = np.array(results.mean_test_score).reshape(6, 6)

mglearn.tools.heatmap(scores, xlabel='gamma', xticklabels=param_grid['gamma'],
                        ylabel='C', yticklabels=param_grid['C'], cmap="viridis")
```

```
Out[31]: <matplotlib.collections.PolyCollection at 0x237f78219a0>
```



The color encodes the cross-validation accuracy, with light colors meaning high accuracy and dark colors meaning low accuracy. You can see that SVC is very sensitive to the setting of the parameters.

Confusion Matrix

```
In [32]: from sklearn.datasets import load_breast_cancer
         from sklearn.metrics import confusion_matrix
```

```
In [33]: cancer = load_breast_cancer()

X_train, X_test, y_train, y_test = train_test_split(
    cancer.data, cancer.target, random_state=0)

logreg = LogisticRegression(solver='liblinear').fit(X_train, y_train)
```


```
pred_logreg = logreg.predict(X_test)
print("logreg score: {:.2f}".format(logreg.score(X_test, y_test)))
```

logreg score: 0.96

```
In [34]: confusion = confusion_matrix(y_test, pred_logreg)
print("Confusion matrix:\n{}".format(confusion))
```

Confusion matrix:
[[52 1]
[5 85]]

```
In [35]: mglearn.plots.plot_binary_confusion_matrix()
```



$$\text{Accuracy} = \frac{TP+TN}{TP+TN+FP+FN} \quad \text{Precision} = \frac{TP}{TP+FP} \quad \text{Recall} = \frac{TP}{TP+FN}$$

$$F = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

```
In [36]: from sklearn.metrics import f1_score

print("f1 score logistic regression: {:.2f}".format(
    f1_score(y_test, pred_logreg)))
```

f1 score logistic regression: 0.97

Threshold

1. Precision-recall curves
2. ROC curves

```
In [37]: #PRC
from sklearn.metrics import precision_recall_curve
```

```
precision, recall, thresholds = precision_recall_curve(  
    y_test, pred_logreg)
```

```
In [38]: # Find the best threshold (e.g., maximize precision and recall)  
optimal_idx = np.argmax(precision + recall) # Choose threshold that gives highest prec  
optimal_threshold = thresholds[optimal_idx]
```

```
In [39]: optimized_predictions = (pred_logreg >= optimal_threshold).astype(int)
```

```
In [40]: print("f1 score logistic regression: {:.2f}".format(  
    f1_score(y_test, optimized_predictions)))
```

f1 score logistic regression: 0.97

```
In [41]: #ROC  
from sklearn.metrics import roc_curve  
fpr, tpr, thresholds = roc_curve(y_test, pred_logreg)
```


```
In [42]: probs = logreg.predict_proba(X_test)[: , 1]
```

```
In [43]: from sklearn.metrics import roc_auc_score  
log_auc = roc_auc_score(pred_logreg, probs)  
  
print("AUC for Logistics: {:.3f}".format(log_auc))
```

AUC for Logistics: 1.000

```
In [44]: plt.plot(fpr, tpr, label="ROC Curve")  
plt.xlabel("FPR")  
plt.ylabel("TPR (recall)")  
# find threshold closest to zero  
close_zero = np.argmin(np.abs(thresholds))  
plt.plot(fpr[close_zero], tpr[close_zero], 'o', markersize=10,  
    label="threshold zero", fillstyle="none", c='k', mew=2)  
plt.legend(loc=4)
```

Out[44]: <matplotlib.legend.Legend at 0x237f76df1f0>



Metrics for Multiclass Classification

```
In [45]: from sklearn.metrics import accuracy_score
```

```
In [47]: from sklearn.datasets import load_digits
digits = load_digits()
```

```
In [48]: X_train, X_test, y_train, y_test = train_test_split(
        digits.data, digits.target, random_state=0)
```

```
In [49]: lr = LogisticRegression().fit(X_train, y_train)
pred = lr.predict(X_test)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

```
In [50]: print("Accuracy: {:.3f}".format(accuracy_score(y_test, pred)))
        print("Confusion matrix:\n{}".format(confusion_matrix(y_test, pred)))
```

Accuracy: 0.951

Confusion matrix:

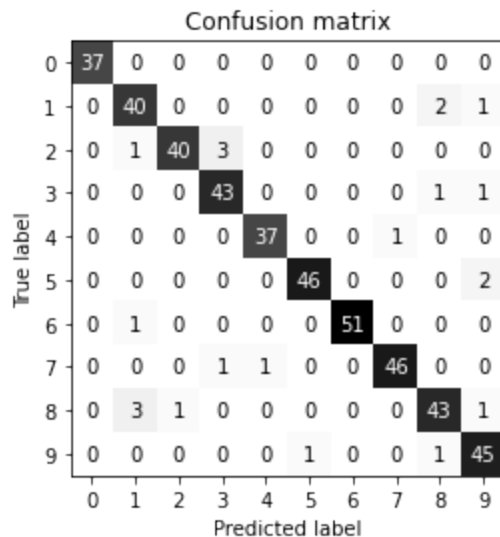
```
[[37  0  0  0  0  0  0  0  0  0]
 [ 0 40  0  0  0  0  0  0  2  1]
 [ 0  1 40  3  0  0  0  0  0  0]
 [ 0  0  0 43  0  0  0  0  1  1]
 [ 0  0  0  0 37  0  0  1  0  0]
 [ 0  0  0  0  0 46  0  0  0  2]
 [ 0  1  0  0  0  0 51  0  0  0]
 [ 0  0  0  1  1  0  0 46  0  0]
 [ 0  3  1  0  0  0  0  0 43  1]
 [ 0  0  0  0  0  1  0  0  1 45]]
```

The model has an accuracy of 95.3%.

In [53]:

```
scores_image = mglearn.tools.heatmap(
    confusion_matrix(y_test, pred), xlabel='Predicted label',
    ylabel='True label', xticklabels=digits.target_names,
    yticklabels=digits.target_names, cmap=plt.cm.gray_r, fmt="%d")

plt.title("Confusion matrix")
plt.gca().invert_yaxis()
```



In [56]:

```
from sklearn.metrics import classification_report

print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	37
1	0.89	0.93	0.91	43
2	0.98	0.91	0.94	44
3	0.91	0.96	0.93	45
4	0.97	0.97	0.97	38
5	0.98	0.96	0.97	48
6	1.00	0.98	0.99	52
7	0.98	0.96	0.97	48
8	0.91	0.90	0.91	48
9	0.90	0.96	0.93	47
accuracy			0.95	450

Model_Evaluation				
macro avg	0.95	0.95	0.95	450
weighted avg	0.95	0.95	0.95	450